

STANDARD CLASS

Running Gemma 4 31B: locally and in the cloud

One 96 GB card locally, one H100 in the cloud. The model every FlatClaw tenant runs today. **Model:** Gemma 4 31B-IT · 31B dense · Apache 2.0 · 256K native · ~33 GB at FP8 + KV cache.

96 GB

VRAM, ONE CARD

FP8NATIVE ON BLACKWELL /
HOPPER**256K**

CONTEXT SERVED

8-12

CONCURRENT STREAMS

1

CARD, ONE PROCESS

LOCAL · BOUGHT ONCE

\$20,235

The one-card inference workstation as specced below; Newegg street prices on October 7, 2026.

CLOUD · HELD WARM

\$2,000/mo

GPU node, 1 × H100 80 GB · 26 vCPU · 234 GB RAM on the reference plan; indicative list price, 24/7, on Azure, AWS, Google Cloud or our reference lane.

01 Local — The one-card inference workstation

One 96 GB card holds Gemma 4 31B at FP8 with the full 256K window, exactly as the hosted tenant runs it; the rest of the box is an ordinary high-end desktop platform. Every part is sold by Newegg today at the price shown.

PART	PICK · WHY IT'S HERE	STREET	BUY
GPU	NVIDIA RTX PRO 6000 Blackwell Max-Q, 96 GB 96 GB of GDDR7 holds the 31B at FP8 with the full 256K window on one card, exactly how the hosted tenant runs it. 300 W, a quiet blower, display outputs. The 600 W Workstation Edition is a few percent faster when Newegg has it in stock; today it does not. The passive Server Edition is wrong for a tower.	\$17,000 in stock	Newegg ↗
CPU	AMD Ryzen 9 9900X, 12 cores Feeds the card over PCIe 5.0 and runs the OS, SGLang, the gateway and ingestion. Twelve cores is plenty; the 9950X only pays off for CPU-heavy ingestion.	\$379 in stock · list \$403	Newegg ↗
MAINBOARD	ASUS ProArt X870E-Creator WiFi Two PCIe 5.0 x16 slots (x8/x8) so a second card can follow later, 10 GbE + 2.5 GbE, four M.2, USB4: workstation-grade AM5.	\$512 in stock	Newegg ↗
MEMORY	Corsair Vengeance 64 GB (2 × 32 GB) DDR5-6000 CL30 Two DIMMs keep AM5 at full speed; 64 GB carries the OS, Docker, the gateway and ingestion with headroom, because the model lives in VRAM. DDR5 has roughly quadrupled in price this year; 96 GB kits are not worth it for this box.	\$1,030 in stock · list \$1,069	Newegg ↗
POWER	Corsair HX1200i (2025), 1200 W, 80+ Platinum, ATX 3.1 A native 12V-2×6 lead for the Pro 6000 and wide margin over the ~500 W peak of card plus platform; the same unit carries the 600 W Workstation Edition if you swap in later.	\$288 in stock · list \$345	Newegg ↗
CPU COOLER	Corsair Nautilus II 360 RS, 360 mm AIO Holds the 120 W 9900X while the GPU dumps heat into the case; fits the North XL's top mount.	\$110 in stock	Newegg ↗
CASE	Fractal Design North XL, Mesh, charcoal and walnut A tower that looks like furniture: 413 mm of GPU room, a 360 mm top radiator, a mesh side feeding the blower. Mesh, not glass.	\$166 in stock · list \$195	Newegg ↗
STORAGE	Samsung 9100 PRO 4 TB, NVMe Gen5 OS plus the model library (~1 TB system, ~3 TB weights); the board takes Gen5 natively.	\$750 in stock · list \$1,360	Newegg ↗

All-in, as specced · 8 parts, 8 priced live from newegg.com on October 7, 2026**\$20,235**

Newegg street prices as read on October 7, 2026; the live table at flatclaw.org/builds refreshes hourly. Used-market and unlisted parts are marked. Verify at checkout.

02 Cloud — the same checkpoint on the cloud you already run

The same checkpoint in a tenancy on the cloud the customer already runs: Azure, AWS, Google Cloud, or our reference lane.

RECOMMENDED · GPU NODE

\$2,000/mo · \$2.74/hr

1 × H100 80 GB · 26 vCPU · 234 GB RAM on the reference plan. About \$2,000 a month held warm 24/7 including the bundled CPU and RAM. Weights are fetched from the weights volume over the tenancy's own network at boot; nothing leaves the tenancy.

Azure NC H100 v5 · AWS p5 · Google Cloud A3 · Northflank nf-gpu-h100-80-1g

PRICE BASIS

Indicative list prices held warm 24/7 on our reference lane; the same classes on Azure, AWS and Google Cloud land in the same band on reserved terms, and bare metal amortizes lower.

NOTES

- ▶ Flat per-tenant rate, not per-token metering: the GPU line dominates and everything else is under \$200.
- ▶ Capacity is peak concurrent sessions, not head count: one H100 sustains 8–12 concurrent streaming chats at FP8 with first tokens in 1–2 s.

CONTROL PLANE

\$144/mo

4 vCPU · 16 GB · the Portal plus one OpenClaw gateway per user. Two or three users at 0.6–1.2 GB each with room to spare. Every tenant has this service next to its GPU, on any cloud's small VM or container plan.

Azure D4s v5 · AWS m6i.xlarge · Google Cloud n2-standard-4 · Northflank nf-compute-400-16

WEIGHTS VOLUME

\$30/mo

200 GB of storage plus a small CPU pod that serves it over HTTP. Staged once; the GPU node fetches the weights at boot because GPU plans rarely mount a volume directly.

Azure Blob + blobfuse2 · AWS S3 · Google Cloud GCS · Northflank volume + CPU pod

ALL-IN

About \$2,200 a month per tenant held warm: the GPU node, the control plane and the weights volume, on whichever cloud the customer already runs.

03 Run it

Serve

```
python -m sglang.launch_server --model-path google/gemma-4-31b-it --quantization fp8 --tp 1 --context-length 262144 --tool-call-parser gemma4 --reasoning-parser gemma4 --port 30000 --mem-fraction-static 0.90
```

- ▶ Peak draw is the card's 300 W plus about 200 W of platform: ~500 W, about 4 A at 120 V on an ordinary circuit. Quiet enough for a desk.
- ▶ Ubuntu 24.04 Server with NVIDIA R570+ / CUDA 12.8+ (sm_120). The card has display outputs for a monitor when you want one and runs headless otherwise. Serve from containers and hold driver updates before demo days.
- ▶ One SGLang process on the one card behind one OpenAI-compatible endpoint, the same launch line as the hosted tenant. The model never reloads.
- ▶ Run the OpenClaw gateway and the Portal on the same box for a fully local setup, or point a separate control machine at its endpoint; both work.